DeepSeek QA

Utrecht University

February 27, 2025

Universiteit Utrecht

Overview

- Training recipe
- Notable details
- Relevance & Discussion

Universiteit Utrecht

DeepSeek Training recipe

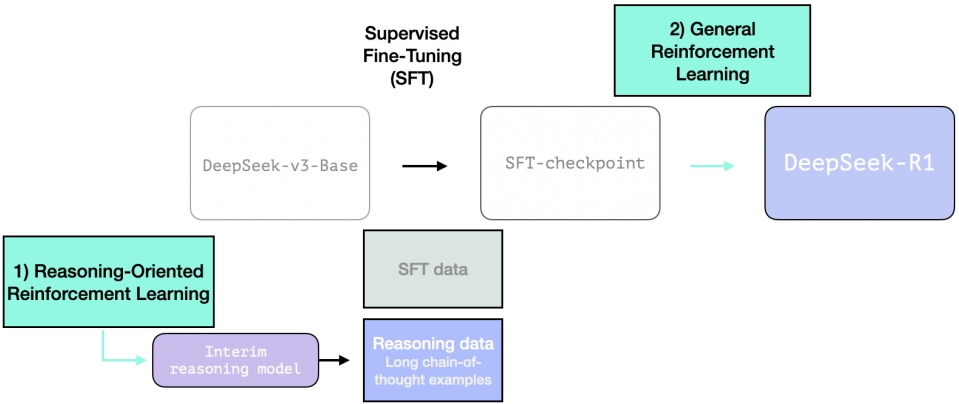

Training: step 1 & 2

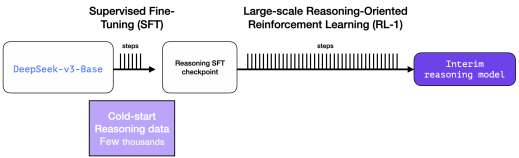

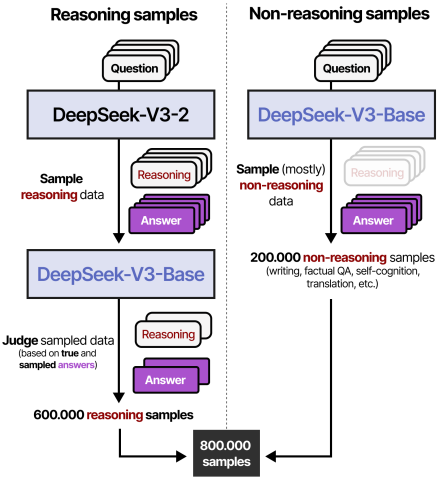

RL process

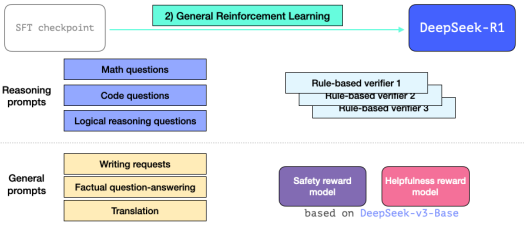
SFT checkpoint
2) General Reinforcement Learning
DeepSeek-R1
Reasoning prompts
Math questions
Code questions
Logical reasoning questions
Rule-based verifier 1
Rule-based verifier 2
Rule-based verifier 3
General prompts
Writing requests
Factual question-answering
Translation
Safety reward model
Helpfulness reward model
based on DeepSeek-v3-Base

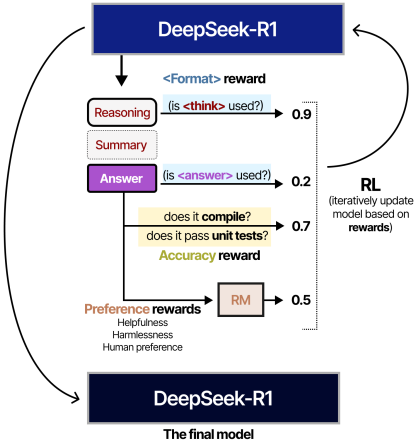
DeepSeek-R1
<Format> reward
Reasoning
(is <think> used?)
0.9
Summary
Answer
(is <answer> used?)
0.2
RL
(iteratively update model based on rewards)
does it compile?
does it pass unit tests?
0.7
Accuracy reward
Preference rewards
RM
0.5
Helpfulness
Harmlessness
Human preference
DeepSeek-R1
The final model

Details: Data?

What does the RL data look like?

1. Unreleased by DeepSeek
2. Reproduced by Open-R1 (and openly released).
3. https://huggingface.co/datasets/open-r1/OpenThoughts-114k-math

Details: GRPO

1. Avoiding the challenge of learning a value function from a LM backbone, where research hasn't established best practices.
2. Saves memory by not needing to keep another set of model weights in memory.

GRPO does this by simplifying the value estimation and assigning the same value to every token in the episode (i.e. in the completion to a prompt, each token gets assigned the same value rather than discounted rewards in a standard value function) by estimating the advantage or baseline. The estimate is done by collecting multiple completions (a_i) and rewards (r_i), i.e. a Monte Carlo estimate, from the same initial state / prompt (s).

To state this formally, the GRPO objective is very similar to the PPO objective above:

$$J(\theta) = \frac{1}{G}\sum_{i=1}^{G}\left(\min\left(\frac{\pi_\theta(a_i|s)}{\pi_{\theta_{old}}(a_i|s)}A_i, \text{clip}\left(\frac{\pi_\theta(a_i|s)}{\pi_{\theta_{old}}(a_i|s)}, 1-\varepsilon, 1+\varepsilon\right)A_i\right) - \beta D_{KL}(\pi_\theta||\pi_{ref})\right).$$

Note that relative to PPO, the standard implementation of GRPO includes the KL distance in the loss. With the advantage computation for the completion index i:

$$A_i = \frac{r_i - \text{mean}(r_1, r_2, \cdots, r_G)}{\text{std}(r_1, r_2, \cdots, r_G)}. \qquad (4)$$

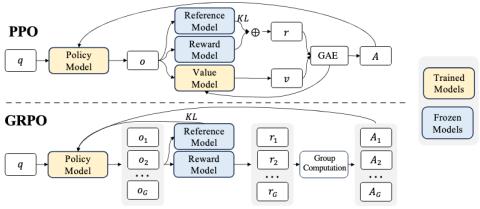

Figure 4 | Demonstration of PPO and our GRPO. GRPO foregoes the value model, instead estimating the baseline from group scores, significantly reducing training resources.

Universiteit Utrecht

Relevance

Can our projects be relevant in this paradigm?

1. Downsampling SFT and reasoning data = data pruning?
2. GRPO: alternative distance functions to KL divergence.
3. Alternatives to coding as reasoning tasks?

Universiteit Utrecht

Discussion

References:

1. https://arxiv.org/abs/2412.19437
2. https://newsletter.languagemodels.co/p/the-illustrated-deepseek-r1
3. https://www.interconnects.ai/p/deepseek-r1-recipe-for-o1
4. https://newsletter.maartengrootendorst.com/p/a-visual-guide-to-reasoning-llms